

PyTorch_NN

March 6, 2018

PyTorch Neural Networks

0.1 Define the network

Neural networks can be constructed using the `torch.nn` package.

`nn` depends on `autograd` to define models and differentiate them. An `nn.Module` contains layers, and a method `forward(input)` that returns the output.

A typical training procedure for a neural network is as follows:

- Define the neural network that has some learnable parameters (or weights)
- Iterate over a dataset of inputs
- Process input through the network
- Compute the loss (how far is the output from being correct)
- Propagate gradients back into the network's parameters
- Update the weights of the network, typically using a simple update rule: $\text{weight} = \text{weight} - \text{learning_rate} * \text{gradient}$

Here is an example of a simple feed-forward network. It has following layers: 1.conv, 2. sub-sampling, 3. conv, 4. subsampling, 5. full connection, 6. full connection, 7. Gaussian connections.

```
In [1]: import torch
        from torch.autograd import Variable
        import torch.nn as nn
        import torch.nn.functional as F

        class Net(nn.Module):

            def __init__(self):
                super(Net, self).__init__()
                # 1 input image channel, 6 output channels, 5x5 square convolution
                # kernel
                self.conv1 = nn.Conv2d(1, 6, 5)
                self.conv2 = nn.Conv2d(6, 16, 5)
                # an affine operation: y = Wx + b
                self.fc1 = nn.Linear(16 * 5 * 5, 120)
                self.fc2 = nn.Linear(120, 84)
                self.fc3 = nn.Linear(84, 10)
```

```

def forward(self, x):
    # Max pooling over a (2, 2) window
    x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
    # If the size is a square you can only specify a single number
    x = F.max_pool2d(F.relu(self.conv2(x)), 2)
    x = x.view(-1, self.num_flat_features(x))
    x = F.relu(self.fc1(x))
    x = F.relu(self.fc2(x))
    x = self.fc3(x)
    return x

def num_flat_features(self, x):
    size = x.size()[1:] # all dimensions except the batch dimension
    num_features = 1
    for s in size:
        num_features *= s
    return num_features

net = Net()
print(net)

Net (
  (conv1): Conv2d(1, 6, kernel_size=(5, 5), stride=(1, 1))
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear (400 -> 120)
  (fc2): Linear (120 -> 84)
  (fc3): Linear (84 -> 10)
)

```

You just have to define the forward function, and the backward function (where gradients are computed) is automatically defined for you using autograd. You can use any of the Tensor operations in the forward function.

The learnable parameters of a model are returned by `net.parameters()`.

```

In [2]: params = list(net.parameters()) # Use net.parameters() to see the parameters.
        print(len(params))
        for i in range(len(params)):
            print(params[i].size())

10
torch.Size([6, 1, 5, 5])
torch.Size([6])
torch.Size([16, 6, 5, 5])
torch.Size([16])
torch.Size([120, 400])
torch.Size([120])

```

```
torch.Size([84, 120])
torch.Size([84])
torch.Size([10, 84])
torch.Size([10])
```

```
In [3]: input = Variable(torch.randn(1, 1, 32, 32))
        out = net(input) # this object is "callable", look at explanations below.
        print(out)
```

Variable containing:

```
0.0210  0.0804 -0.0869  0.0386  0.0267  0.0725 -0.1889 -0.0007 -0.1421  0.1407
[torch.FloatTensor of size 1x10]
```

`forward()` is a method in `nn.Module` class. Document said that it should be overridden by all subclasses.

Callable

Look at the `net(input)` statement (highlighted by a "*"). This object is *callable*. For the definition of "callable" objects, here is a short introduction from Stack Overflow. "The built-in callable (PyCallable_Check in objects.c) checks if the argument is either:

- an instance of a class with a **call** method or
- is of a type that has a non null `tp_call` (c struct) member which indicates callability otherwise (such as in functions, methods etc.)

The method named `__call__` is called when the instance is "called" as a function.

Since the `Net()` class is extended from `nn.Module`, each of its subclasses should override the `forward()` method. Actually, the `__call__` of `nn.Module` calls the `forward()`, which is overridden in each of the subclasses.

```
In [4]: net.zero_grad() # Zero the gradient buffers of all parameters and backprops with random
        out.backward(torch.randn(1,10)) # the argument in Variable.backward() is gradient.
```

`torch.nn` only supports **mini-batches** The entire `torch.nn` package only supports inputs that are a mini-batch of samples, and **not** a single sample. For example, `nn.Conv2d` will take in a 4D Tensor of `nSamples x nChannels x Height x Width`. If you have a single sample, just use `input.unsqueeze(0)` to add a fake batch dimension.

More Recap: - `nn.Parameter`: A kind of Variable, that is automatically registered as a parameter when assigned as an attribute to a Module. - `autograd.Function`: Implements forward and backward definitions of an autograd operation. Every Variable operation, creates at least a single Function node, that connects to functions that created a Variable and encodes its history.

0.2 Loss function

A loss function takes the (output, target) pair of inputs, and computes a value that estimates how far away the output is from the target. There are several different loss functions under the `nn` package . A simple loss is: `nn.MSELoss` which computes the mean-squared error between the input and the target. Use *Dash* to search `torch.nn` for more information.

```
In [5]: output = net(input)
        target = Variable(torch.arange(1,11))
        print("Target Data:")
        print(target.data)
        print("Output Data:")
        print()
```

Target Data:

```
1
2
3
4
5
6
7
8
9
10
```

[torch.FloatTensor of size 10]

Output Data:

```
In [6]: criterion = nn.MSELoss()
        loss = criterion(output, target)
        print(loss)
```

Variable containing:

```
38.6209
```

[torch.FloatTensor of size 1]

```
In [7]: print(loss.grad_fn)
        print(loss.grad_fn.next_functions[0][0])
        print(loss.grad_fn.next_functions[0][0]
              .next_functions[0][0])
```

```
<torch.autograd.function.MSELossBackward object at 0x7fce17cb5408>
```

```
<torch.autograd.function.AddmmBackward object at 0x7fce17cb5318>
```

```
<AccumulateGrad object at 0x7fce64777438>
```

0.3 Backprop

To backpropagate the error all we have to do is to `loss.backward()`. You need to clear the existing gradients though, else gradients will be accumulated to existing gradients. Now we shall call `loss.backward()`, and have a look at `conv1`'s bias gradients before and after the backward.

```
In [8]: net.zero_grad()
        print('conv1.bias.grad before backward')
        print(net.conv1.bias.grad)
```

conv1.bias.grad before backward

Variable containing:

```
0
0
0
0
0
0
0
```

[torch.FloatTensor of size 6]

```
In [9]: loss.backward()

        print('conv1.bias.grad after backward')
        print(net.conv1.bias.grad)
```

conv1.bias.grad after backward

Variable containing:

```
0.0847
0.1343
-0.0375
-0.0675
0.1161
0.0870
```

[torch.FloatTensor of size 6]

0.4 Update the weights

The simplest update rule used in practice is the Stochastic Gradient Descent (SGD):

$$\text{weight} = \text{weight} - \text{learning_rate} * \text{gradient}$$

We can implement this using simple python code:

```
learning_rate = 0.01
for f in net.parameters():
    f.data.sub_(f.grad.data * learning_rate)
```

However, as you use neural networks, you want to use various different update rules such as SGD, Nesterov-SGD, Adam, RMSProp, etc. To enable this, we can use a small package: `torch.optim` that implements all these methods. Here's how to use it:

```
In [10]: import torch.optim as optim
```

```
# Create your optimizer
optimizer = optim.SGD(net.parameters(), lr=0.01)

# in your training loop
optimizer.zero_grad() # zero the gradient buffers
output = net(input)
loss = criterion(output, target)
loss.backward()
optimizer.step() # Does the update
```